

Winch — Quote Recovery Opportunity Assessment

Version 2.1 · 18 September 2026 · Supersedes "Evaluating Trade Quote Recovery Opportunity" v1.0

Target markets: United Kingdom · Europe · Australia · wider EMEA

v2.1 change. The customer-facing channel is no longer SMS. It is ****WhatsApp with email fallback**** — see section 5 and [docs/DESIGN.md](#). SMS was dropped to avoid a paid third party and because WhatsApp's approved-template buttons give the customer a one-tap reply, which suits this audience better than composing a text. Everything else below stands.

>

What changed in v2.0. v1.0 claimed that existing CRMs do not serve this market. That claim is false and has been removed — Jobber, ServiceTitan and ServiceM8 all ship quote follow-up automation today. The thesis has been re-grounded on **non-adoption** rather than absence. The "35% of high-ticket quotes die" figure has been demoted from a fact to an unvalidated hypothesis. Geography has been narrowed from "universal regional problem" to UK/EU/AU/EMEA, because the WhatsApp control plane does not work in the US or Canada. The human-in-the-loop design has been revised to survive an unresponsive contractor.

1. The product in one line

A LangGraph-powered agent that lives inside WhatsApp, ingests a contractor's existing quote PDFs, and runs an approved, human-gated follow-up sequence to the end customer over plain-text SMS — so that high-ticket quotes stop dying in silence.

Positioned to the buyer as **"the virtual estimating assistant that lives in your WhatsApp"**, never as "an AI agent app".

2. The market

Regional and peri-urban trades outside the metro core: shed and barn builders, commercial and agricultural fencing, borehole and water-well drilling, heavy equipment mechanics, agricultural fabricators, groundworks and civils subcontractors.

Attribute	Profile
Job value	GBP 8k–60k / EUR 10k–70k / AUD 15k–110k per quote
Business size	Owner-operator to ~10 field staff; no inside sales, no estimator
Quote volume	Roughly 8–30 quotes per month
Who quotes	The owner, typically evenings and Sundays, after a full day in the field
Who follows up	Nobody. There is no role that owns it.

The buying psychology that governs everything

- **Reputation-protective.** They will not tolerate anything that reads as pushy, automated or "sleazy" going out under their name. This is the single largest adoption constraint and the reason strict human-in-the-loop is a feature, not an overhead.
- **The technology paradox.** They will spend six figures on an excavator or a CNC machine without hesitation, then agonise over a EUR 150/month subscription. Capital equipment is tangible; software is not.
- **Zero tolerance for setup.** If it takes more than about fifteen minutes to understand, they revert to a spreadsheet and text messages.
- **Channel reality.** In the UK, Ireland, most of Europe, Australia and the wider EMEA region, the owner already runs their working life through WhatsApp — crews, suppliers, site photos, everything. The app is already open.

3. Pain points — assessed

Each pain point below is graded on whether it survives contact with evidence, not on whether it sounds compelling.

Pain point	Verdict	What is actually true
Dormant quotes never followed up	Real — the core thesis	Structural, not a preference. A one-truck owner working 12-hour days has no role that owns "day 5 check-in". The work has no home, so it does not happen. This does not self-correct.
"The software is too complex to navigate"	Real, but misdiagnosed	The blocker is not UI complexity — Jobber's mobile app is genuinely good and onboards inside 48 hours. The blocker is desk time . Every CRM assumes someone sits down to do data entry, configure triggers and maintain a pipeline. A field owner never sits down.
Quote data entry friction	Real — strongest agent fit	The owner already has a PDF. Forwarding it into a chat thread is near-zero friction, and structured extraction is exactly what this technology is good at.
Customers live on text, not email	Real	Site managers, farm owners and self-builders answer a text from a local number. They do not open a corporate-looking email template five days after receiving a quote.
Offline / low-connectivity working	Not a real problem for us	WhatsApp already queues and sends on reconnect. This only becomes a problem if we build a native mobile app — so we will not build one.
"Existing CRMs fail this market"	False — removed from the thesis	See section 4.

The 35% figure

v1.0 asserted that up to 35% of high-ticket quotes die purely from lack of follow-up. That number is unsourced and has the shape of vendor marketing. **Do not put it in a deck or a pitch.** Treat it as the primary hypothesis to be measured, and measure it in the first pilot — baseline quote-to-close rate before, quote-to-close rate after. That data exists by week three of a pilot and is worth more than any borrowed statistic.

4. The correction that changes the pitch

Existing CRMs do **not** fail this market on features. Jobber ships automated quote reminders on its Grow tier. ServiceTitan runs drip campaigns against open quotes. ServiceM8 already markets AI-assisted quoting to tradies. The capability exists, and some of it is cheaper than our intended price point.

The real finding is stronger and more defensible than the false one:

We are not competing with absence. We are competing with non-adoption. Those follow-up features sit switched off and unconfigured, because configuring them is precisely the desk work the owner will not do. The gap is not a missing feature. It is a missing *zero-onboarding surface* for a feature that already exists.

This has three consequences:

- 1 **Never pitch "nobody solves this."** A prospect with Jobber on their phone will catch it, and you lose the room. Pitch instead: *"You already pay for follow-up automation. Let's find out if it's switched on."* It is true, it is disarming, and it is usually a yes.
- 2 **Our moat is distribution and interface, not capability.** The defensible asset is that the contractor does nothing but forward a PDF and reply to a message in an app they already have open.
- 3 **Integrations become a second-order concern.** Xero (UK/AU), QuickBooks, Sage and regional invoicing tools matter eventually, but the wedge works on forwarded PDFs alone. Do not gate the first pilot on an accounting integration.

5. Architecture — the dual-sided model

The most common failure in conversational automation is treating the contractor and the end customer as two ends of one chatbot. They are two different systems with different regulatory constraints, different tone, and different levels of autonomy.

Side A — Contractor to agent: WhatsApp control plane

The contractor never logs into a web app. They forward a quote PDF, photograph a handwritten work order, or send a ten-second voice note into a dedicated WhatsApp Business number. The agent confirms the extracted data, proposes a cadence, and requests approval via interactive buttons or quick replies. It pings them only when a customer responds with real buying intent.

Because the contractor always initiates, this side stays comfortably inside Meta's 24-hour service window. That is not a workaround — it is why this side belongs on WhatsApp and the other does not.

Side B — Agent to customer: WhatsApp primary, email fallback

Customer-facing outreach goes over WhatsApp where the customer is reachable there, and over email where they are not. No SMS, and therefore no paid third-party messaging provider.

The customer's number and email are **collected from the contractor** at confirmation rather than parsed from the quote, because most trade quote PDFs contain no mobile number at all.

There is no way to pre-check whether a number is on WhatsApp — Meta removed that endpoint and keeps delivery errors deliberately vague. Resolution is therefore send-and-observe: attempt WhatsApp, and on a 131026 error or no delivery confirmation inside the timeout, fall back to email permanently for that quote.

Because the five-day follow-up is business-initiated and outside the 24-hour window, the cold touch must be a **pre-approved template**. Button labels are frozen at approval time and capped at three, so the agent selects from a small approved catalogue rather than generating buttons per message. The moment the customer taps a button or replies, the 24-hour window opens and free-form interactive messages become available.

Email sends as the contractor, through their own Gmail on OAuth. It threads onto the original quote email, arrives from their real address, and has no template approval, no 24-hour window and no per-message fee — the fallback is the less constrained channel.

Regulatory note — Meta platform terms

Meta's Business Solution terms, effective 15 January 2026, bar general-purpose AI assistants from WhatsApp. Task-specific business bots — support, orders, bookings, and our case, quote management — remain permitted. We sit on the right side of this line, but the contractor-facing agent must stay visibly scoped to quotes and follow-ups and must not be allowed to drift into a general-purpose chat assistant. This is a product constraint, not just a compliance note.

6. Why LangGraph rather than Zapier or Make

Quoting is non-linear. A contractor revises a quote mid-sequence. A customer pays a partial deposit. A deal goes dormant for three weeks waiting on planning permission or a council permit, then wakes up. Linear automation tools break on all three.

Primitive	Use in the recovery engine	Why it is necessary
Multimodal extraction sub-agent	Ingests arbitrary PDFs, photos of work orders, messy spreadsheets	Normalises inconsistent regional proposals into one strict schema: <code>customer_name</code> , <code>customer_phone</code> , <code>project_title</code> , <code>quote_total</code> , <code>expiry_date</code>
<code>interrupt()</code> HITL gate	Pauses graph execution before any external message	The contractor's reputation is the asset at risk. Nothing leaves the system unapproved.
Checkpointed persistence	Long-horizon state across weeks	Quote lifecycles run 2 to 30+ days. The graph must survive restarts and remember touchpoint history and objections.
Triage sub-agent	Classifies inbound customer replies	Routes to <code>QUESTION_ON_TIMELINE</code> , <code>PRICE_OBJECTION</code> , <code>ACCEPTED</code> , <code>TECHNICAL_SCOPE_QUERY</code> , <code>UNSUBSCRIBE</code>

7. Hard engineering rules

These are non-negotiable and belong in the codebase, not just the document.

- 1 The LLM never generates or alters a monetary figure, line item, or delivery date.** Quote totals and scope travel through state as immutable literals. A hallucinated price texted to a customer is contractual and financial exposure for the contractor, and reputational death for us.
- 2 Every outbound customer message passes a human gate.** No exceptions, including scheduled touchpoints. See section 8 for how the gate is designed so that it does not stall.
- 3 The customer never converses with an unconstrained model.** Any inbound reply touching scope, price or technical tolerance halts the sequence immediately and hands the contractor a tap-to-call link.
- 4 Timing is deterministic, never agentic.** Wait states, cadences and cron intervals live in Postgres or a Redis queue with LangGraph checkpoints. The model does extraction, tone and classification. It does not decide to wait three days.
- 5 STOP and unsubscribe handling ships before the first real customer message.** Non-negotiable in the UK and EU, and the fastest way to destroy a contractor's reputation if it fails.

- 6 **Data protection by default.** UK GDPR and EU GDPR apply to end-customer contact data from day one: lawful basis recorded, retention limits set, processor terms in place with the contractor. Build it in now; retrofitting it across EMEA is expensive.

8. The design risk that could sink the product

The diagnosed pain is *"the owner never does the follow-up."* A strict gate-per-message design requires the owner to tap approve on **every single follow-up**. That is the same fatigue, re-pointed at a new inbox. An owner who ignores an unanswered quote will equally ignore a 6pm approval prompt — and the sequence stalls silently while they pay a monthly retainer for it.

Strict human-in-the-loop remains correct for trust and legal exposure. The fix is to change *when* the human decides, not *whether*.

Front-loaded approval with a veto window

- **Approve once, at ingestion.** When the contractor forwards the PDF they are already in the thread and engaged. That is the moment of maximum attention. The agent presents the parsed quote and the full proposed sequence, and asks for one approval covering all three touchpoints.
- **Veto, do not gate, on each send.** Two hours before each scheduled touchpoint the agent posts a preview: *"Sending Mark the day-5 schedule nudge at 4pm — reply HOLD to stop it."* Silence means send. Control is preserved; attention is not required.
- **Hard gate only on exceptions.** A customer reply, a scope question, a price objection or anything outside the approved sequence stops everything and waits for a human. These are rare, high-signal, and worth interrupting for.

This preserves 100% contractor control over content while surviving an owner who goes quiet for a week — which is the normal case, not the edge case.

Treat this as the core design decision of the demo, not an implementation detail. It is also the second thing to validate with real contractors (see the companion validation plan).

9. Commercial model

Element	Position
Pricing	GBP 250–350 / EUR 300–400 / AUD 500–650 per month, plus a setup fee. Anchor on "one recovered job pays for the year."
Messaging cost	Near-zero. WhatsApp template sends run roughly GBP 0.02 each; email is free. No SMS provider.
Recovered-revenue share	Attractive in theory, deferred in practice. Proving a deal was "dead" is an attribution fight you will lose in month two. Start flat; revisit once you have baseline data.
Proof of ROI	Reviving one EUR 25k shed build or engine rebuild every two months returns several times the retainer. The maths is easy; the <i>evidence</i> is the hard part.
Churn risk	Quote recovery is lumpy. A quiet month with no saves feels like paying for nothing. Mitigate with a monthly one-line summary: quotes tracked, touchpoints sent, replies generated — visible activity between wins.

Route to market

- **Do not buy search or LinkedIn ads.** Returns are poor against this persona.
- **Channel partnerships are the unlock.** Regional steel and timber suppliers, equipment dealers, and trade associations. If a major steel supplier recommends the tool to its commercial fencing accounts to help move inventory, you inherit trust that advertising cannot buy. This works especially well across the UK, Ireland, Benelux and Australia, where regional supplier relationships are dense and long-standing.
- **The single-fold demo.** The entire sales motion is one sentence: *"Show me a quote you sent last month that went cold. Let me run a three-step text sequence on it, free. If they reply, we talk."*

10. Geography

Region	Verdict	Note
UK & Ireland	Primary	WhatsApp-native, dense regional trade networks, Xero and Sage penetration, clear GDPR regime
Australia	Primary	WhatsApp-native among owners, high-value rural trades, Xero dominance, existing tradie-software culture (ServiceM8, Tradify)
EU (Benelux, DACH, Iberia, Nordics, CEE)	Secondary, staged	Strong WhatsApp penetration; fragmented invoicing tools and language localisation are the cost
Wider EMEA (Middle East, South Africa)	Opportunistic	WhatsApp is the default business channel; procurement and payment norms differ enough to need separate work
US & Canada	Excluded	Trades live on SMS and phone calls. WhatsApp reads as a foreign app. The control-plane wedge does not exist here — a different product would be required.

Pick one primary market and do not hedge. The messaging norms, the invoicing stack and the partner networks are all country-specific, and splitting effort early halves the quality of the evidence you collect.

11. Verdict

The pain is real. The channel wedge is real. The "software gap" framing was not, and has been removed.

Build it — but the demo's job is not to prove that an agent can parse a PDF and send a text. Anyone can demo that. The demo's job is to show a contractor **their own cold quote from last month receiving a reply**. That is the only proof this buyer responds to, and it is achievable in a 48-hour build.

Two things must be validated with real contractors before significant engineering:

- 1 **Do they already own follow-up automation and simply not use it?** This confirms or kills the non-adoption thesis on which the entire positioning now rests.
- 2 **Will they actually respond to an approval prompt at 9pm on a Sunday?** This confirms or kills strict per-message human-in-the-loop, and determines the core interaction design.

See the companion document: **Validation Plan — 5 Contractor Interviews (UK / EU / AU / EMEA)**.